

# Spark Project Readiness Assessment (Beginner → Deployment Gauge)

**Audience:** Trainees with basic exposure to Python/Spark.

**Goal:** Decide if a beginner can be safely deployed to real projects.

**What we assess:** Foundations, Python coding efficiency, Spark fundamentals & architecture, hands-on data processing, troubleshooting, and delivery hygiene.

---

## 1) Outcomes & Decision Bands

- **Deploy Now** ( $\geq 80/100$  and all gates passed): Can own low/medium-complexity Spark tasks with mentorship for design reviews.
- **Deploy with Mentorship** (65–79/100, gates passed): Pair on critical paths; independent on well-scoped tickets.
- **Not Yet** ( $< 65/100$  or any gate failed): Remediation plan before deployment.

**Hard Gates (must pass):** 1. Environment set up end-to-end (local + cluster access) and reproducible. 2. Code compiles, runs, and is version-controlled with a clear README & configs. 3. Spark job completes on the baseline dataset without errors/OOM.

---

## 2) Assessment Flow (Step-by-Step)

**Duration:** ~1–2 days total hands-on time.

### 1. Stage 0 – Environment & Hygiene (Gate)

2. Install: Python 3.10+, Spark 3.x/4.x, Java 11/17, IDE, Git.
3. Prove: `spark-shell` / `pyspark` start; `spark-submit` runs a hello-world job.
4. Repo layout, virtual env, `.env` /config separation, logging enabled, unit test scaffold.

### 5. Stage 1 – Python Foundations (Core Skills)

6. Write 3 small utilities (functions/classes) + simple unit tests.
7. Data wrangling exercise using pandas or pure Python for CSV/JSON.

### 8. Stage 2 – Spark Fundamentals & Architecture (Knowledge Check)

9. Short quiz + whiteboard: DAG → stages → tasks, shuffle, partitions, joins, Catalyst/Tungsten, broadcast/cache, Spark UI, deploy modes.

### 10. Stage 3 – Spark Core + SQL (Hands-On)

11. Build a small ETL: read → clean → transform (joins, windows) → write partitioned Parquet.
12. Show Spark UI screenshots; explain plan; justify optimizations.

### 13. Stage 4 – Troubleshooting Lab

14. Fix an intentionally broken job (schema drift, skew, nulls, small files, OOM). Document root cause + fix.

### 15. Stage 5 – Mini-Project Capstone

16. End-to-end pipeline with configs, tests, and a `spark-submit` command. Short retro write-up.
- 

## 3) Scoring Weights (100 pts)

- Stage 0: Environment & Hygiene – **Gate only** (no points; fail → stop)
- Stage 1: Python Foundations – **20 pts**
- Stage 2: Spark Fundamentals & Architecture – **20 pts**
- Stage 3: Spark Core + SQL (Hands-On) – **35 pts**
- Stage 4: Troubleshooting Lab – **15 pts**
- Stage 5: Mini-Project Delivery & Hygiene – **10 pts**

**Readiness Bands:** Deploy Now  $\geq 80$ ; Deploy w/ Mentorship 65–79; Not Yet  $< 65$ .

**Bonus (up to +5):** Clean code comments, meaningful docstrings, thoughtful performance notes.

---

## 4) Metrics & Rubrics

### 4.1 Python Foundations (20 pts)

**Tasks** - Implement: `parse_config(path) -> dict` (supports JSON/YAML, env overrides). - Implement: `dedupe_records(records, keys)` preserving original order. - Implement: simple quality checks (null %, unique keys, value ranges) with unit tests.

**Metrics - Correctness (8 pts):** All provided unit tests pass; edge cases handled (empty, nulls, bad types). - **Complexity & Efficiency (4 pts):** Avoid  $O(n^2)$  where  $O(n \log n)$  is feasible; space usage reasonable. - **Readability & Style (4 pts):** PEP8, naming, small functions, type hints. - **Testing (4 pts):**  $\geq 80\%$  coverage on utility module; clear test names; arrange-act-assert format.

**Thresholds** - Must pass all reference tests;  $< 5$  lint warnings per 200 LOC.

---

### 4.2 Spark Fundamentals & Architecture (20 pts)

**Format:** 15–20 short questions + 2 whiteboard prompts.

**What good looks like** - Explains: driver vs executor, jobs→stages→tasks, narrow vs wide transformations. - Can read Spark UI: identify shuffle stages, skewed tasks, spilled bytes. - Knows when to use: broadcast join, cache/persist levels, coalesce vs repartition. - Understands: Catalyst optimizer (logical→physical plan), Tungsten/codegen, AQE basics.

**Rubric - Concept Accuracy (12 pts):** Definitions & trade-offs correct. - **Diagnosis Reasoning (4 pts):** Given a DAG/UI snippet, states likely bottleneck. - **Architecture Mapping (4 pts):** Maps deploy mode (client/cluster) and cluster manager (YARN/K8s/Standalone) to where driver/executors live.

**Thresholds** -  $\geq 70\%$  quiz; both whiteboards coherent (not necessarily perfect).

---

### 4.3 Spark Core + SQL – Hands-On (35 pts)

**Dataset (provide to trainee):**

- `transactions` (~5–10M rows synthetic): id, user\_id, ts, amount, status, merchant, category.
- `users_dim` (~100k): user\_id, join\_date, city, plan.

**Required Pipeline** 1. Ingest CSV/JSON with schema (no schema-on-read), bad records to quarantine. 2. Clean: parse timestamps, trim strings, handle nulls, enforce types. 3. Transform: - Join `transactions` ↔ `users_dim` (broadcast if small). - Compute 7/30-day rolling sums per user (window). - Flag anomalies (amount > p99 per user plan). 4. Write: partitioned Parquet by `dt=YYYY-MM-DD`, overwrite dynamic partitions.

**Deliverables** - Code (PySpark or Scala), `application.conf`, `logging.conf`, `run_local.sh`, `spark-submit` cmd. - Query plan (`df.explain("extended")`) and 2 Spark UI screenshots (stages, SQL tab).

**Metrics - Correctness (12 pts):** Row counts per step, null policy enforced, dedupe rules applied. - **Performance (10 pts):** - Uses predicate/file pruning, avoids unnecessary UDFs. - Broadcast join when `users_dim < 100MB` OR justified alternative. - Shuffle stages minimized; skew handled if present. - **Resource Use (5 pts):** No OOM; spill reduced; sensible `spark.sql.shuffle.partitions`. - **Code Quality (5 pts):** Modular jobs, params via config/args; logs key metrics. - **Reproducibility (3 pts):** One-command run; deterministic outputs.

**Thresholds** - Pipeline completes within **2×** assessor baseline on provided machine/cluster. -  $\leq 3$  shuffle stages unless justified; job succeeds on 2× input scale.

---

### 4.4 Troubleshooting Lab (15 pts)

**Broken Scenarios (assessor provides any 3):** 1. **Schema drift:** unexpected column, wrong type → job fails.

2. **Data skew:** a few keys dominate → stage stragglers.
3. **Small files:** 1000 tiny partitions in lake → slow reads.
4. **Out-of-memory:** collect/action misuse; huge shuffle; wrong cache level.
5. **Bad joins:** exploding row counts due to duplicate keys.

**Candidate Tasks** - Identify root cause using logs + Spark UI.

- Implement fix (e.g., `broadcast`, salting, `coalesce/repartition`, combine small files, map-side pre-agg, filter early, checkpointing).

**Rubric - Diagnosis (6 pts):** Reads UI metrics (task time, input/shuffle read, spill).

- **Fix Quality (6 pts):** Correct and minimally invasive; measures before/after.

- **Communication (3 pts):** Clear incident note with cause→change→result.

**Thresholds** - Fix  $\geq 2$  of 3 scenarios to passing quality; each shows measurable improvement.

---

## 4.5 Mini-Project Delivery & Hygiene (10 pts)

**Requirements** - `README` with problem statement, data contract, run steps, config matrix.

- `tests/` with at least 5 unit tests for pure functions and 2 job-level assertions.

- Parameterized job (env, date range).

- Outputs: partitioned Parquet + data quality report (row counts, null %, distinct keys, p95/p99).

**Rubric - Repo Hygiene (3 pts):** Clear structure, `.gitignore`, versioned configs; no secrets.

- **Testing & CI stub (3 pts):** Local test run + placeholder CI script.

- **Observability (2 pts):** Logs with run id, input/output counts, duration.

- **Docs (2 pts):** Assumptions, limitations, and next steps.

---

## 5) Student Expectation Guide (share with trainees)

**What we'll measure:** - You can write clean Python utilities with tests. - You understand Spark's execution model and can read the Spark UI. - You can build an ETL that joins, aggregates (window), and writes partitioned data efficiently. - You can diagnose/fix common Spark issues (skew, small files, schema drift, OOM). - Your work is reproducible, documented, and version-controlled.

**Coding Standards** - Prefer built-ins & Spark SQL functions over UDFs; if UDFs, justify and unit-test.

- Avoid `collect()` / `show()` on large data; sample with `limit()` / `sample()`. - Log input/output counts, partitions, and plan hints at INFO level. - Parameterize paths, dates, and thresholds via config/args; no hard-coded credentials.

**Performance Principles** - Prune early (columns/rows).

- Broadcast small dims; filter before join; aggregate before shuffle.

- Coalesce writes to avoid tiny files; partition on query-friendly columns.

- Cache only when reused; use correct storage levels; unpersist.

---

## 6) Evaluation Templates

### 6.1 Scoring Sheet (per candidate)

Area	Max	Score	Notes
Python Foundations	20		
Spark Fundamentals & Architecture	20		
Spark Core + SQL (Hands-On)	35		

Area	Max	Score	Notes
Troubleshooting Lab	15		
Mini-Project Delivery & Hygiene	10		
<b>Total</b>	<b>100</b>		
Decision (Deploy Now / Deploy w/ Mentorship / Not Yet)			

## 6.2 Run Log (attach to PR/README)

- Dataset size (rows/GB):
- Cluster/machine spec:
- `spark-submit` args:
- Stage metrics (from UI): input rows, shuffle read/write, spill, skewed tasks.
- Job duration (first run / optimized run):
- Notes & follow-ups:

## 6.3 Incident Note (for Troubleshooting)

- **Symptom:**
- **Root Cause:**
- **Fix Applied:**
- **Evidence:** (plan/UI metrics before/after)
- **Residual Risk:**

## 7) Assessor Pack (what you provide)

- Baseline dataset generator or CSV dumps (transactions + users).
- Reference config + baseline `spark-submit` (with partitions).
- Broken job variants for the lab (toggle via config flags).
- Answer key for fundamentals quiz (not shared with trainees until after).

## 8) Remediation Plan (if Not Yet)

- **Week 1:** Python drills (functions, tests, file IO), re-attempt Stage 1.
- **Week 2:** Spark foundations deep-dive + Spark UI reading workshop, re-quiz.
- **Week 3:** Guided ETL rebuild with performance checklist, re-attempt Stage 3.
- **Exit:** Redo Troubleshooting Lab with new scenarios.

## 9) Quick Reference (Glossary)

- **Driver/Executor:** Orchestration vs work processes.
- **Job/Stage/Task:** Execution hierarchy; stage = shuffle boundary.
- **Narrow/Wide Transformations:** Whether shuffle occurs.
- **Broadcast Join:** Send small table to all executors.
- **AQE:** Adapts joins/partitions at runtime.

- **Skew:** Uneven key distribution causing stragglers.
  - **Small Files:** Many tiny output files harm read performance.
- 

## Appendix A – Suggested Repo Skeleton

```
project-root/  
  src/  
    main/python/etl/  
      __init__.py  
    jobs/  
      daily_enrich.py  
    utils/  
      io.py  
      quality.py  
      config.py  
  tests/  
    test_quality.py  
  conf/  
    application.conf  
    logging.conf  
  scripts/  
    run_local.sh  
  README.md  
  requirements.txt (or pyproject.toml)  
  .gitignore
```

## Appendix B – Sample `spark-submit`

```
spark-submit \  
  --master local[*] \  
  --deploy-mode client \  
  --conf spark.sql.shuffle.partitions=200 \  
  --conf spark.sql.adaptive.enabled=true \  
  --conf spark.sql.broadcastTimeout=300 \  
  src/main/python/etl/jobs/daily_enrich.py \  
  --date 2025-08-01 --env local --config conf/application.conf
```

---

Use this document as both the candidate handout (expectations) and the assessor checklist (rubrics + scoring).